

NeuroSymbolic Commonsense Reasoner

Project 1: Text-Based Reasoning Module

CS-F429 Course Project | November 25, 2025

Introduction & Core Objectives

Project Goal

To build a **NeuroSymbolic pipeline** capable of commonsense reasoning. The system is designed to bridge the gap between the flexibility of neural models and the rigor of symbolic logic.

Task 1 Approach

We focus exclusively on text-based reasoning using the **CLUTRR dataset**. Our strategy combines LLM-based semantic parsing (Mistral-7B) with deterministic symbolic inference (Prolog) to ensure accuracy and explainability.

Problem Statement

Convert natural-language family stories into structured predicates, perform multi-hop logical inference, and generate a final relationship with a Chain-of-Thought explanation.

The CLUTRR Dataset

id string · lengths	story string · lengths	query string · lengths	target int32	target_text string · classes 18 values
36	36	45	769	14
0fc660c1-e7d5-41fb-8d72-...	[Ashley]'s daughter,...	('Ashley', 'Nicholas')	15	son
b1c985cf-69b5-4575-97d3-...	[Nancy] likes to cut the hair of...	('Nancy', 'Lorraine')	16	daughter
5e7fc867-3782-4ff6-98d5-...	[Dale] and his sister [Nancy] ar...	('Dale', 'Louise')	17	niece
9bc318aa-3c85-4fef-8e6f-...	[Lillian] and her sister [Nancy] ar...	('Nancy', 'Douglas')	14	nephew
6cb675ce-a3d1-44ca-8d3e-...	[Ashley] liked to go to the park...	('Dale', 'Ashley')	6	mother
689c12a6-d83d-4ad4-9bf0...	[Theresa] wants to make a special...	('Ashley', 'Dale')	11	granddaughter

Compositional Language Understanding and Text-based Relational Reasoning


- Content: Short stories describing complex family relationships.
- Complexity: Requires 2 to 4 hops of reasoning to solve.
- Objective: Identify the specific relationship between two target entities within the narrative.
- Role: Serves as the ground truth for testing our neuro-symbolic extraction capabilities.

Architecture Overview




1. Semantic Extraction

Using Mistral-7B Instruct to parse natural language into structured logic predicates.



2. Symbolic Reasoning

Executing multi-hop inference using a local SWI-Prolog solver for deterministic results.



3. CoT Generation

Verbalizing the Prolog solver traces back into human-readable Chain-of-Thought explanations.

Stage 1: Semantic Extraction

Using Mistral-7B Instruct

We employ a frozen Mistral-7B model to extract logical predicates from the story text. To ensure compatibility with our solver, we enforced strict constraints:

Prompt Engineering Strategy:

- Closed vocabulary of relations (e.g., **father**, **aunt**).
- Python-style list output format.
- Two-pass COT pipeline: Summary → Predicates.

Query extraction:

- A simple multi-shot prompt with suitable examples was sufficient.

```
actual story : [Lynn] loves cooking for her son. His name is [Thomas] [Thomas] wanted to go to his grandfather [James] 'house because he always has a good time when he is there. [James] bought a new dress for his daughter [Lena].
```

```
Running model1.....
```

```
The following generation flags are not valid and may be ignored: ['temperature']. Set 'TRANSFORMERS_VERBOSITY=info' for more details.
```

1. Lynn is Thomas's mother.
2. Thomas wants to visit James.
3. James is Lena's father.
4. Lena received a new dress from James.

```
Running model2.....
```

```
The following generation flags are not valid and may be ignored: ['temperature']. Set 'TRANSFORMERS_VERBOSITY=info' for more details.
```

```
sections:
```

```
Genders:  
lynn=female.  
thomas=male.  
james=male.  
lena=female.
```

```
-----  
Triples:  
(thomas, mother, lynn),  
(james, father, lena)
```

```
-----  
Facts:  
father(james, lena).  
mother(lynn, thomas).  
male(thomas).  
female(lynn).  
male(james).  
female(lena).
```

Extraction Challenges & Solutions

Hallucinations & Formatting: LLMs often produce verbose text or hallucinated relationships outside the schema.

Solution: Extract valid subject-object pairs and filter against ALLOWED_INPUTS.

Missing Predicates: If a relationship type (e.g., "uncle") isn't in the story, the Prolog solver crashes on undefined predicates.

Solution: Automatic dummy fact injection (e.g., uncle(dummy1, dummy2)) ensures the Knowledge Base is complete.

Syntax Drift: Ensuring the output is strictly valid Prolog code.

Solution: Strict few-shot prompting with explicit negative constraints (e.g., "No conversational text") and a two step pipeline.

LLM Stability: Trying to maintain only the required context in the LLM.

Solution: Two step pipeline (Story -> Summary -> Predicates)

Stage-1.1 Prompts

```
p1 = """You are a family-relationship extraction engine.

Your job is to read the story and rewrite it into relationship bullet points that contains:
- only and all family relationships
- gender every new person you come across using context - dont base it off the name
- NO events, objects, activities, locations, or irrelevant descriptions
- NO timeline or sequencing details
- NO pronouns; always replace pronouns with the correct resolved person name
- NO hallucinations – include ONLY what is explicitly stated or directly implied
- If a relationship is ambiguous, OMIT it.

extract relation phrases from the story after excluding the unnecessary details using all previous context :

{story}
FINISH
"""
```

Stage-1.1

```
actual story : [Emma] went to the garden to find her father [John]. [John] was talking to his brother [Peter]. [Peter] mentioned that their mother [Sarah] was inside watching TV. [Emma]'s sister [Kate] ran past them playing tag. Later, [John]'s wife [Lisa] called everyone for lunch.
```

```
Running model1.....
```

```
The following generation flags are not valid and may be ignored: ['temperature']. Set 'TRANSFORMERS_VERBOSITY=info' for more details.
```

1. Emma is the daughter of John.
2. John is the father of Emma and the brother of Peter.
3. Peter is the brother of John and the father of Kate.
4. Sarah is the mother of John and Peter.
5. Kate is the sister of Emma and the daughter of Peter.
6. Lisa is the wife of John.

Stage-1.2 Prompts

```
p2 = """You are a family-relationship mapper.  
Your task is to read the cleaned family summary and output ONLY three sections:  
Genders, Triplets, Facts.  
  
You MUST think step-by-step internally, but DO NOT reveal your reasoning.  
Output ONLY the three sections exactly as described.  
  
=====  
RULES  
=====  
1. Names refer to unique people.  
2. Use only explicit relationships from the summary.  
3. Infer gender from roles (father, mother, son, daughter, husband, wife, brother, sister) DO NOT INFER FROM THEIR NAMES!.  
4. DO NOT guess information that is not stated.  
5. DO NOT create placeholder names or "unknown".  
6. Use the most specific relation available.  
7. In the Facts section, include gender as:  
    male(name). female(name). unknown_gender(name).  
(Use lowercase names as atoms.)  
  
=====  
OUTPUT SPECIFICATION  
=====  
  
### SECTION 1 – Genders  
Format: name=male/female/unknown  
  
### SECTION 2 – Triplets  
Format: (subject, relation, object), relations are only family relations, ignore others:  
ONLY Allowed relation labels:  
('granddaughter', 'father', 'brother', 'mother', 'grandson', 'sister', 'uncle', 'daughter', 'husband', 'grandmother', 'wife', 'aunt', 'son', 'grandfather')  
  
### SECTION 3 – Facts  
Convert both gender and triplets into Prolog facts.  
  
For gender:  
- male(name).  
- female(name).  
- unknown_gender(name).  
  
For relations:  
father(subject, object),  
brother(subject, object), etc  
  
Use lowercase atoms:  
Example: father(don, steve).  
  
=====  
NOW PROCESS THESE RELATIONS into sections:  
{story}  
RELATIONS FINISH  
sections:  
"""
```

Stage-1.2

```
Running model2.....  
  
sections:  
-----  
  
Genders:  
Emma=female,  
John=male,  
Peter=male,  
Sarah=female,  
Kate=female,  
Lisa=unknown  
  
Triplets:  
(John, father, Emma),  
(John, brother, Peter),  
(Peter, father, Kate),  
(Peter, brother, John),  
(Sarah, mother, John),  
(Sarah, mother, Peter),  
(Kate, sister, Emma),  
(Kate, daughter, Peter)  
  
Facts:  
male(john).  
female(emma).  
female(sarah).  
male(peter).  
female(kate).  
unknown_gender(lisa).  
father(john, emma).  
brother(john, peter).  
father(peter, kate).  
brother(peter, john).  
mother(sarah, john).  
mother(sarah, peter).  
sister(kate, emma).  
daughter(kate, peter)
```

Stage-1.3

```
def to_prolog_query(question: str):
    prompt = f"""
Convert the following natural-language question into a Prolog query.
Use ONLY standard Prolog syntax.

RULES:
- Use lowercase atoms for predicates and constants.
- Replace spaces in names with underscores.
- For yes/no questions, generate a predicate that logically matches the question.
- Output ONLY the Prolog query, nothing else.

EXAMPLES:
EXAMPLES:

Q: "Is Alice Bob's mother?"
A: mother(alice, bob).

Q: "Is Bob Alice's father?"
A: father(bob, alice).

Q: "Is John Mary's brother?"
A: brother(john, mary).

Q: "Is Sarah Emily's sister?"
A: sister(sarah, emily).

Q: "Is Claire the grandmother of David?"
A: grandmother(claire, david).

Q: "Is Helen David's granddaughter?"
A: granddaughter(helen, david).

Q: "Is Michael Jane's son?"
A: son(michael, jane).

Q: "Is Jane Michael's daughter?"
A: daughter(jane, michael).

NOW CONVERT:
Q: "{question}"
A:
"""

    return generate(prompt, max_new_tokens=50, temperature=0.001)
```

```
query = to_prolog_query("is helena chloe's mother in law")
print(query.split("A:")[-1])
```

The following generation flags are not valid and may be ignored: ['temperature']. Set 'TRANSFORMERS_VERBOSITY=info' for more details.

```
mother_in_law(helena, chloe).
```

Stage 2: Symbolic Reasoning

We chose **Prolog** over Z3 or MiniZinc for three key reasons:

Trace Generation: Prolog natively supports SLD

1. resolution, providing a clear proof trace essential for CoT.
2. **Logic Fit:** Perfectly handles the multi-hop recursive logic required for family trees.

Note: None of Z3, MiniZinc, PySwip, PyDatalog, SWI-Prolog were both lightweight enough to run on notebooks and had native trace support (Support was removed for most versions and we faced dependency issues while downgrading)

The sanitized predicates are written to a .pl file alongside a static ruleset (rules.pl).



SWI Prolog

The screenshot shows the SWI-Prolog IDE interface. On the left, the main window displays a trace of a query. The trace starts with a call to `husband/2` with arguments `(13) husband(_5338, gina) ? creep`. It then branches through various predicates like `is_married/2`, `parent/2`, and `father/2` to reach a final answer `X = arthur .`. On the right, there is a smaller window titled "Prolog.Rules.pl - Notepad" containing a set of Prolog rules for basic inference, including gender determination and parentage relations.

```
File Edit Settings Run Debug Help
File Edit Format View Help
% SECTION 2: CORE INFERENCE LOGIC
% =====
% These rules extract "Basic Truths" (Gender, Parenthood, Marriage)
% from the mixed bag of facts above.

% --- Gender Inference ---
% Determine if someone is Male based on provided facts
gender(X, male) :- father(X, _).
gender(X, male) :- brother(X, _).
gender(X, male) :- grandson(X, _).
gender(X, male) :- uncle(X, _).
gender(X, male) :- husband(X, _).
gender(X, male) :- wife(_, X).
gender(X, male) :- son(X, _).
gender(X, male) :- grandfather(X, _).


% Determine if someone is Female based on provided facts
gender(X, female) :- mother(X, _).
gender(X, female) :- sister(X, _).
gender(X, female) :- granddaughter(X, _).
gender(X, female) :- aunt(X, _).
gender(X, female) :- wife(X, _).
gender(X, female) :- husband(_, X).
gender(X, female) :- daughter(X, _).
gender(X, female) :- grandmother(X, _).

% --- Parent Inference ---
% Determine Parent(P, C) relationship from various facts
parent(P, C) :- father(P, C).
parent(P, C) :- mother(P, C).  

Ln 32, Col 23 100% Windows (CRLF) UTF-8
```

Inference Logic

- father
- mother
- brother
- sister
- grandson
- granddaughter
- uncle
- aunt
- husband
- wife
- grandmother
- grandfather
- son
- daughter
- mother-in-law
- father-in-law



Inference Logic

A. The Simple Combinations (Direct Lookup)

These rules simply apply a "Gender Filter" to a "Relationship Link".

- `is_father(X, Y)`

$$= \text{gender}(X, \text{male}) + \text{parent}(X, Y)$$

(Logic: A male who is a parent.)

% 1. FATHER: A male parent
`is_father(X, Y) :-`
 `gender(X, male),`
 `parent(X, Y).`

B. The "Triangulated" Combinations (Siblings)

These rules are smarter. Since you don't have a "sibling" core truth, you construct it by finding a **shared parent**.

- `is_sister(X, Y)`

$$= \text{gender}(X, \text{female}) + \text{parent}(Z, X) + \text{parent}(Z, Y)$$

% 8. SISTER: Female sibling (shares at least one parent)
`is_sister(X, Y) :-`
 `gender(X, female),`
 `parent(Z, X),`
 `parent(Z, Y),`
 `X \neq Y.`

Inference Logic

C. The "Chained" Combinations (Grandparents)

These rules work by "hopping" over an intermediate person.

- `is_grandfather(X, Y)`

$$= \text{gender}(X, \text{male}) + \text{parent}(X, Z) + \text{parent}(Z, Y)$$

(Logic: A male who is the parent of someone (Z), who is the parent of Y.)

```
% 9. GRANDFATHER: Male parent of a parent  
is_grandfather(X, Y) :-  
    gender(X, male),  
    parent(X, Z),  
    parent(Z, Y).
```

```
is_grandfather(X, Z) :-  
    gender(X, male),  
    grandfather(X, Y),  
    brother(Y, Z).
```

```
is_grandfather(X, Z) :-  
    gender(X, male),  
    grandfather(X, Y),  
    sister(Y, Z).
```

D. The "Hybrid" Combinations (In-Laws)

This is your most complex rule type. It combines a **Spouse Link** with a **Parent Link**.

- `is_mother_in_law(X, Y)`

$$= \text{gender}(X, \text{female}) + \text{spouses}(Y, Z) + \text{parent}(X, Z)$$

(Logic: A female who is the parent of the person (Z) that Y is married to.)

```
% 16. MOTHER-IN-LAW: Mother of one's spouse  
is_mother_in_law(X, Y) :-  
    gender(X, female), % X must be female  
    spouses(Y, Z), % Y is married to Z  
    parent(X, Z). % X is the parent of Z (the spouse)
```

```

[trace] ?- is_grandfather(james, john).
Call: (10) is_grandfather(james, john) ? creep
Call: (11) gender(james, male) ? creep
Call: (12) father(james, _2952) ? creep
Exit: (12) father(james, robert) ? creep
Exit: (11) gender(james, male) ? creep
Call: (11) parent(james, _5308) ? creep
Call: (12) father(james, _5308) ? creep
Exit: (12) father(james, robert) ? creep
Exit: (11) parent(james, robert) ? creep
Call: (11) parent(robert, john) ? creep
Call: (12) father(robert, john) ? creep
Fail: (12) father(robert, john) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) mother(robert, john) ? creep
Fail: (12) mother(robert, john) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) son(john, robert) ? creep
Fail: (12) son(john, robert) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) father(robert, _19080) ? creep
Fail: (12) father(robert, _19080) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) father(robert, _21512) ? creep
Fail: (12) father(robert, _21512) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) mother(robert, _23944) ? creep
Fail: (12) mother(robert, _23944) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) mother(robert, _26376) ? creep
Fail: (12) mother(robert, _26376) ? creep
Redo: (11) parent(robert, john) ? creep
Call: (12) mother(james, _5308) ? creep
Fail: (12) mother(james, _5308) ? creep
Redo: (11) parent(james, _62) ? creep
Call: (12) son(_62, james) ? creep
Fail: (12) son(_62, james) ? creep
Redo: (11) parent(james, _62) ? creep
Call: (12) daughter(_62, james) ? creep
Fail: (12) daughter(_62, james) ? creep
Redo: (11) parent(james, _62) ? creep
Call: (12) father(james, _5564) ? creep
Exit: (12) father(james, robert) ? creep
Call: (12) brother(_62, robert) ? creep
Exit: (12) brother(william, robert) ? creep
Exit: (11) parent(james, william) ? creep
Call: (11) parent(william, john) ? creep
Call: (12) father(william, john) ? creep
Fail: (12) father(william, john) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) mother(william, john) ? creep
Fail: (12) mother(william, john) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) son(john, william) ? creep
Fail: (12) son(john, william) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) daughter(john, william) ? creep
Fail: (12) daughter(john, william) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) father(william, _20146) ? creep
Fail: (12) father(william, _20146) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) son(john, william) ? creep
Fail: (12) son(john, william) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) daughter(john, william) ? creep
Fail: (12) daughter(john, william) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) father(william, _22578) ? creep
Fail: (12) father(william, _22578) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) mother(william, _25010) ? creep
Fail: (12) mother(william, _25010) ? creep
Redo: (11) parent(william, john) ? creep
Call: (12) mother(william, _27442) ? creep
Fail: (12) mother(william, _27442) ? creep
Redo: (11) parent(william, john) ? creep
Fail: (11) parent(james, _62) ? creep
Call: (12) father(james, _30684) ? creep
Exit: (12) father(james, robert) ? creep
Call: (12) sister(_62, robert) ? creep
Exit: (12) sister(mary, robert) ? creep
Exit: (11) parent(james, mary) ? creep
Call: (11) parent(mary, john) ? creep
Call: (12) father(mary, john) ? creep
Fail: (12) father(mary, john) ? creep
Redo: (11) parent(mary, john) ? creep
Call: (12) mother(mary, john) ? creep
Fail: (12) mother(mary, john) ? creep
Redo: (11) parent(mary, john) ? creep
Call: (12) son(john, mary) ? creep
Exit: (12) son(john, mary) ? creep
Exit: (11) parent(mary, john) ? creep
Exit: (10) is_grandfather(james, john) ? creep
true .

```



▼	...	Exit: (12) father(james, robert) ? creep Exit: (11) gender(james, male) ? creep Exit: (12) father(james, robert) ? creep Exit: (11) parent(james, robert) ? creep Exit: (12) father(james, robert) ? creep Exit: (12) brother(william, robert) ? creep Exit: (11) parent(james, william) ? creep Exit: (12) father(james, robert) ? creep Exit: (12) sister(mary, robert) ? creep Exit: (11) parent(james, mary) ? creep Exit: (12) son(john, mary) ? creep Exit: (11) parent(mary, john) ? creep Exit: (10) is_grandfather(james, john) ? creep
---	-----	--

Stage 3: Chain-of-Thought (CoT)



Prolog Trace

The solver outputs a raw logical trace
of the derivation steps



Verbalization

Mistral-7B takes the trace and
translates the logical steps into
natural language sentences with COT.



Final Output

A human-friendly explanation (CoT)
accompanying the final relationship
answer.

Stage-3 Prompts

```
def summarize_prolog_trace_cot(trace: str):
    prompt = """
Summarize the following exit-only Prolog trace. First, think step-by-step
about what each exit line implies logically. Then, after reasoning, produce
a clean final paragraph that explains the reasoning chain in natural language.

RULES:
- Use lowercase atoms as they appear.
- DO NOT output the reasoning steps in the final answer.
- Output ONLY the final summarized paragraph.
- Your reasoning should interpret each exit line in order.

EXAMPLES:

TRACE:
Exit: mother(alice, beth)
Exit: parent(beth, claire)
Exit: grandmother(alice, claire)

THINK:
- mother(alice, beth) means Alice is Beth's mother.
- parent(beth, claire) means Beth is Claire's parent.
- Combining them shows Alice is the grandmother of Claire through Beth.

SUMMARY:
Alice is Claire's grandmother because she is the mother of Beth, and Beth is Claire's parent, forming a direct maternal lineage.

---


TRACE:
Exit: father(john, mark)
Exit: parent(mark, olivia)
Exit: grandfather(john, olivia)

THINK:
- father(john, mark) → John is Mark's father.
- parent(mark, olivia) → Mark is Olivia's parent.
- Combine → John is Olivia's grandfather through Mark.

SUMMARY:
John is Olivia's grandfather since he is Mark's father, and Mark is Olivia's parent, creating a paternal line connecting John and Olivia.
```

```
TRACE:
Exit: sibling(sarah, michael)
Exit: parent(michael, emma)
Exit: aunt(sarah, emma)

THINK:
- Sarah and Michael are siblings.
- Michael is Emma's parent.
- A sibling of a parent is an aunt.

SUMMARY:
Sarah is Emma's aunt because she is Michael's sibling, and Michael is Emma's parent.

---


TRACE:
Exit: parent(linda, jack)
Exit: sibling(linda, robert)
Exit: parent(robert, claire)
Exit: cousin(jack, claire)

THINK:
- Linda is Jack's parent.
- Linda and Robert are siblings.
- Robert is Claire's parent.
- Children of siblings are cousins.

SUMMARY:
Jack and Claire are cousins because their parents, Linda and Robert, are siblings.

---


NOW SOLVE:
TRACE:
{trace}

THINK:
"""
return generate(prompt, max_new_tokens=300, temperature=0.001)
```

Stage-3

```
In [12]:  
import re  
  
def extract_exit_traces(trace: str) -> str:  
    exit_lines = re.findall(r"^\s*Exit:\*", trace, flags=re.MULTILINE)  
    new_traces = "\n".join(exit_lines)  
    return new_traces  
  
new_traces = extract_exit_traces(trace)  
print(new_traces)  
  
Exit: (12) mother(sarah, john) ? creep  
Exit: (11) gender(sarah, female) ? creep  
Exit: (12) mother(sarah, john) ? creep  
Exit: (11) parent(sarah, john) ? creep  
Exit: (12) father(john, emma) ? creep  
Exit: (11) parent(john, emma) ? creep  
Exit: (10) is_grandmother(sarah, emma) ? creep
```



```
In [13]:  
COT = summarize_prolog_trace_cot(new_traces)  
  
The following generation flags are not valid and may be ignored: ['temperature']. Set 'TRANSFORMERS_VERBOSE=info' for more details.
```



```
In [14]:  
print(COT.split("THINK:")[-1])  
  
- Sarah is John's mother.  
- Sarah is female.  
- Sarah is John's parent.  
- John is Emma's father.  
- John is Emma's parent.  
- Combining them shows Sarah is Emma's grandmother.  
  
SUMMARY:  
Sarah is Emma's grandmother because she is John's mother, and John is Emma's parent, forming a direct maternal lineage.
```

Thank you!
Questions?